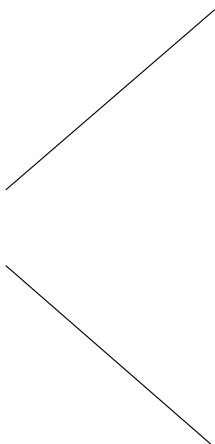


```
$ oc api-resources -o name --sort-by=name
alertmanagers.monitoring.coreos.com
apiservers.config.openshift.io
apiservices.apiregistration.k8s.io
appliedclusterresourcequotas.quota.openshift.io
authentications.config.openshift.io
authentications.operator.openshift.io
baremetalhosts.metal3.io
bindings
brokertemplateinstances.template.openshift.io
buildconfigs.build.openshift.io
builds.build.openshift.io
builds.config.openshift.io
catalogsources.operators.coreos.com
certificatesigningrequests.certificates.k8s.io
cloudcredentials.operator.openshift.io
clusterautoscalers.autoscaling.openshift.io
clusternetworks.network.openshift.io
clusteroperators.config.openshift.io
...
```

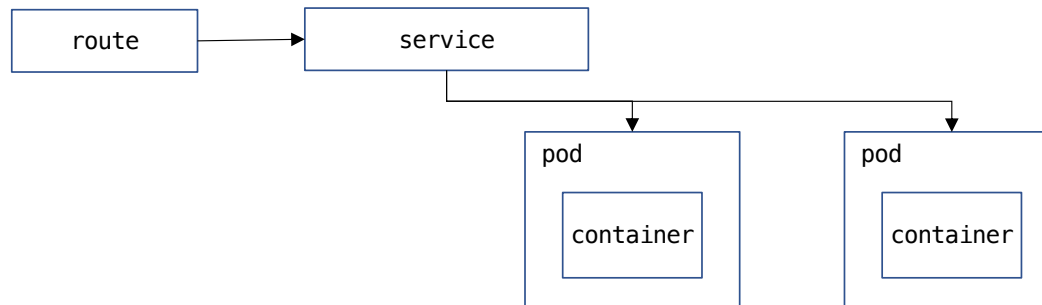


*Pod*  
*Replicationcontroller (rc)*  
*Deploymentconfig (dc)*  
*Service (svc)*  
*Route*  
*PersistenceVolumeClaim (pvc)*  
*Secrets*  
*Configmaps (cm)*  
  
*Imagestream (is)*  
*BuildConfig (bc)*  
  
*Node*  
*PersistenceVolume (pv)*  
  
*Operator*  
*CustomResourceDefinition (crd)*

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wildfly
  labels:
    app: wildfly
spec:
  replicas: 2
  selector:
    matchLabels:
      app: wildfly
  template:
    metadata:
      labels:
        app: wildfly
    spec:
      containers:
        - name: wildfly
          image: quay.io/do288/wildfly:latest
          ports:
            - containerPort: 8080
              protocol: TCP
```

```
apiVersion: v1
kind: Service
metadata:
  name: wildfly
  labels:
    app: wildfly
spec:
  type: ClusterIP
  selector:
    app: wildfly
  ports:
    - name: http
      protocol: TCP
      port: 8080
      targetPort: 8080
```

```
apiVersion: route.openshift.io/v1
kind: Route
metadata:
  name: wildfly
  labels:
    app: wildfly
spec:
  host: sample.apps.eu46.prod.nextcle.com
  to:
    kind: Service
    name: wildfly
  tls:
    termination: edge
```



## Declarative :

```
$ ls
deployment.yml route.yml service.yml

$ oc apply -f .
deployment.apps/wildfly created
route.route.openshift.io/wildfly created
service/wildfly created
```

## Imperative :

```
$ oc new-app <container-image | git-repository>
--> Found container image 9a9e908 (9 days old) from quay.io for "quay.io/do288/wildfly"

    * An image stream tag will be created as "wildfly:latest" that will track this image

--> Creating resources ...
    imagestream.image.openshift.io "wildfly" created
    deployment.apps "wildfly" created
    service "wildfly" created
--> Success
```

```
$ oc new-app --help
Create a new application by specifying source code, templates, and/or images
```

...

Usage:

```
oc new-app (IMAGE | CONTAINERFILE | SOURCE | TEMPLATE | ...) [flags]
```

Beispiele:

```
$ oc new-app quay.io/do288/nginx --name nginx
```

└──────────┘

Container-Image

```
$ oc new-app php:7.3~https://github.com/.../php-hello
```

└──┘ └──────────────────────────┘

Builder-Image  
(s2i)

Git-Projekt (Source)

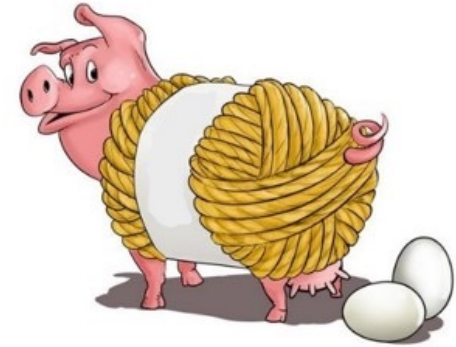


Deployment

Service

Imagestream

BuildConfig



## Imagestream:

- enthält Verweise (Zeiger) auf Images und deren Tags (keine Images)
- Verwendung in Deployment als Image und Trigger oder in BuildConfig als (S2I) BuilderImage
- Import aus externer Registry oder Ergebnis eines Build

```
$ oc import-image nginx --from=quay.io/do288/nginx --confirm ( --scheduled)
$ oc describe is nginx
Name:                               nginx
Unique Images:                      4
Tags:                               4

latest
  updates automatically from registry quay.io/do288/nginx:latest
  * quay.io/do288/nginx@sha256:c34f57431167fca470730b67a1a8636126d2464eee619ec8d0b577c8e63bffe0

1.2
  updates automatically from registry quay.io/do288/nginx:1.2
  * quay.io/do288/nginx@sha256:ee508edacfe0bc1e6af43a15348b400a7d97121507348bd5fb5effb6b9f8d84e

1.1
  updates automatically from registry quay.io/do288/nginx:1.1
  * quay.io/do288/nginx@sha256:674ab485f6e83f162eb4bdaf12986469c7b4f484f65fbb18f3b03218fd5f36e4

1.0
  updates automatically from registry quay.io/do288/nginx:1.0
  * quay.io/do288/nginx@sha256:693b30b107da26
```

TAG	LAST MODIFIED ↓	SECURITY SCAN	SIZE	MANIFEST
1.2	40 minutes ago	8 Medium	91.9 MB	SHA256 ee508edacfe0
latest	14 hours ago	8 Medium	91.9 MB	SHA256 c34f57431167
1.1	a day ago	8 Medium	90.6 MB	SHA256 674ab485f6e8
1.0	a day ago	8 Medium	90.6 MB	SHA256 693b30b107da



- `oc login -u <user> -p <password> <api-server-url>`
- `oc new-project <name>`
- `oc create -f <resource-yml>`  
`oc apply -f <resource-yml>`
- `oc status`
- `oc get <resource-type> [ <resource-name> ]`
  - `oc get pods`
  - `oc get deployment`
  - `oc get svc <service>`
  - `oc get events`
- `oc describe <resource-type> <resource-name>`
- `oc expose svc <service-name>`
- `oc logs <podname>`
- `oc exec -it <podname> -- <program>`
- `oc rsh <podname>`
- `oc cp <pod>:<location> <location>`
- `oc port-forward <podname> <local-port>:<remote-port>`
- `oc new-app <☺anything☺>`
- `oc delete <resource-type> <resource-name>`
- `oc rollout latest deployment <deployment-name>`

[https://docs.openshift.com/container-platform/4.12/cli\\_reference/openshift\\_cli/developer-cli-commands.html](https://docs.openshift.com/container-platform/4.12/cli_reference/openshift_cli/developer-cli-commands.html)

## podman build - Containerfile

```
FROM ubi8/ubi
LABEL version=.... maintainer=
MAINTAINER daniel
ENV key=value
ARG version
```

```
ADD http://repos/app-$version.tar /opt/app/
COPY webapp.war /opt/tomcat/webapps
RUN yum install -y tomcat && \
    useradd tomcat && \
    chgrp -R 0 /opt/tomcat && \
    yum cleanup && rm /tmp/*
```

```
ONBUILD COPY . /tmp/src
USER 1000
WORKDIR /opt/tomcat
VOLUMES /opt/tomcat/logs
EXPOSE [ 8080, 8001 ]
ENTRYPOINT [ "bin/sh -c" ]
CMD [ "bin/catalina.sh", "start" ]
```

```
podman build --format docker -t my-tomcat . <-
Build-Dir
```

```
podman push <registry>/<name>:<tag>
```

Layer

Config

Layer

...

Layer

Config

Manifest

local container storage

remote registry

## Verwenden von YUM/DNF beim Image-Build

```
$ podman run --rm ubi8/ubi cat /etc/yum.repos.d/ubi.repo
[ubi-8-baseos]
name = Red Hat Universal Base Image 8 (RPMs) – BaseOS
baseurl = https://cdn-ubi.redhat.com/content/public/ubi/dist/ubi8/8/$basearch/baseos/os
enabled = 1
gpgkey = file:///etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
gpgcheck = 1
...
```


yum "telefoniert"  
nach aussen !

Lösung: beim podman-build andere yum-Konfiguration (z.B. vom Host) mounten

Bei Verwendung von Satellite/Subscriptions ggf. auch die notwendigen Zertifikate/GPG Schlüssel.

```
$ sudo podman build -v /etc/yum.repos.d:/etc/yum.repos.d -v /etc/pki:/etc/pki -v /etc/rhsm:/etc/rhsm .
```

## Beispiel: Image – Vererbung



Änderungen an einem Basis-Image erfordern Rebuild der davon abhängigen Images !

## Container in Openshift:

- beliebige User-Id      RUN chmod - R 0770 ....
- Group-Id 0 (root)      RUN chgrp -R 0
- Ports > 1024

```
apiVersion: project.openshift.io/v1
kind: Project
metadata:
  annotations:
    openshift.io/sa.scc.mcs: s0:c26,c15
    openshift.io/sa.scc.supplemental-groups: 1000680000/10000
    openshift.io/sa.scc.uid-range: 1000680000/10000
```

```
# oc exec pgadmin-778c479f79-tfbqn -- id
uid=1000680000(1000680000) gid=0(root) groups=0(root),1000680000
```

NFS-Mount →

```
# ls -al /mnt/nfs/apps/pgadmin
-rw-r--r-- 1 1000680000 root 124K Nov 27 01:03 access_log
-rw-r--r-- 1 1000680000 root 853 Nov 27 00:44 config_local.py
-rw-r--r-- 1 1000680000 root 1.2K Nov 27 00:46 error_log
```

<https://cloud.redhat.com/blog/a-guide-to-openshift-and-uids>

Abweichende User-Id : Serviceaccount mit Security Context Constraint 'anyuid' notwendig :

```
apiVersion:
rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: scc-anyuid
rules:
- apiGroups:
  - security.openshift.io
  resourceNames:
  - anyuid
  resources:
  - securitycontextconstraints
  verbs:
  - use
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: gitea:anyuid
  namespace: apps
roleRef:
  kind: ClusterRole
  name: scc-anyuid
  apiGroup: rbac.authorization.k8s.io
subjects:
- kind: ServiceAccount
  name: gitea
  namespace: apps
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: gitea
  namespace: apps
```

erstellt von Cluster-Administrator !

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: gitea
  namespace: apps
...
spec:
  template:
    spec:
      serviceAccountName: gitea
  ...
```

```
# oc exec gitea-7dcdc5c445-w9qmv -- id
uid=65534(nobody) gid=65534(nobody) groups=65534(nobody),0(root)

# ll /mnt/nfs/repos/ds
drwxr-xr-x 7 nobody nobody 119 Nov 26 16:57 admin.git/
drwxr-xr-x 7 nobody nobody 119 Nov 26 16:12 calibre.git/
drwxr-xr-x 7 nobody nobody 119 Nov 17 16:02 gitea.git/
...
```

UserId aus Container-Config !

## Secrets:

- Passwörter, Token, Zertifikate ...
- typisiert: basic-auth, dockerfg, tls, opaque
- Inhalte sind base64-decodiert, nicht verschlüsselt

→ max. Größe 1 MB

→ nur innerhalb eines Project (NS) sichtbar

```
apiVersion: v1
kind: Secret
metadata:
  name: ...
  namespace: ...
data:
  password: MTIzNDU2
type: Opaque
```

```
# echo MTIzNDU2 | base64 -d
123456
```

## ConfigMap:

- generische Key-Value Daten

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: ...
  namespace: ...
binaryData:
  keystore: |
    7oAMCAQICCF7Dt6ZDf6TgMA0GCSqGSIb3DQEBBQUAMEI1ZSQUla
    MTEQMA4GA1UECwwHU ...
data:
  HOME: /usr/share/nginx
  default.conf: |
    server {
      listen 8181 default_server;
      server_name _;
      location / {
        root /usr/share/nginx/html;
        index index.html index.htm;
      }
    }
```

```
$ oc create configmap <cm-name> --from-literal F00=BAR
```

```
$ oc create configmap <cm-name> --from-file <path>
```

```
$ oc create secret docker-registry quayio --docker-server quay.io --docker-username <user> --docker-password <password>
```

## Secrets: Verwendung als Umgebungs-Variabe

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-env-pod
spec:
  containers:
  - name: mycontainer
    image: redis
    env:
    - name: SECRET_USERNAME
      valueFrom:
        secretKeyRef:
          name: mysecret
          key: username
    - name: SECRET_PASSWORD
      valueFrom:
        secretKeyRef:
          name: mysecret
          key: password
```

```
$ oc set env deployment/<deployment-name> --from cm/<cm-name>
```

```
$ oc set volume deployment/<deployment-name> --add --name config --configmap-name <cm-name>
```

## ConfigMap: Verwendung als Konfigurations-Dateien

```
apiVersion: apps/v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    container: nginx
    volumeMounts:
    - mountPath: /etc/nginx/conf.d
      name: config
  volumes:
  - name: config
    configMap:
      name: nginx-config
```

```
apiVersion: apps/v1
kind: Pod
metadata:
  name: wildfly-standalone-xml
spec:
  containers:
  - name: wildfly
    container: nginx
    volumeMounts:
    - mountPath: /opt/wildfly/standalone/configuration
      name: standalone-xml
      subPath: standalone.xml
  volumes:
  - name: standalone-xml
    configMap:
      name: standalone-xml
```



## Container Registry:

Red Hat → <https://access.redhat.com/RegistryAuthentication>

```
# podman login quay.io
Username: ...
Password: ...
Login Succeeded!    -> /run/user/<user-id>/containers/auth.json
```

```
# podman push --creds <username>:<password> ...
```

```
# skopeo --help
Various operations with container images and container image registries
```

Usage:  
 skopeo [command]

### Available Commands:

copy	Copy an IMAGE-NAME from one location to another
delete	Delete image IMAGE-NAME
help	Help about any command
inspect	Inspect image IMAGE-NAME
list-tags	List tags in the transport/repository specified by the REPOSITORY-NAME
login	Login to a container registry
logout	Logout of a container registry
manifest-digest	Compute a manifest digest of a file
standalone-sign	Create a signature using local files
standalone-verify	Verify a signature using local files
sync	Synchronize one or more images from one location to another

```
skopeo copy --format ... --dest-creds <user>:<password> containers-storage:localhost/webserver docker://quay.io/do288/webserver
```

## Verwenden einer externen Container Registry - Authentifizierung

```
$ oc get serviceaccounts
```

```
NAME          SECRETS
```

```
builder       2
```

```
default       2
```

```
deployer      2
```

```
$ oc create secret docker-registry quayio --docker-server quay.io --docker-username <user> --docker-password <password>
```

```
$ oc secrets link default quayio --for pull
```

```
$ oc secrets link builder quayio
```

## Verwenden einer externen Container Registry - Secret von auth.json

```
$ oc create secret generic quayio --from-file .dockerconfigjson=/run/user/1000/containers/auth.json --type kubernetes.io/dockerconfigjson
```

```
apiVersion: v1
kind: Secret
metadata:
  name: quayio
type: kubernetes.io/dockerconfigjson
data:
  .dockerconfigjson: ewogICJhdXRocyI6IHsKICAgICJyZWdp3 ...
```

### Serviceaccount 'imagePullSecrets' :

```
$ oc secrets link default quayio --for pull
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: default
imagePullSecrets:
- name: default-dockercfg-4sdrk
- name: quayio
...
```

oder im Deployment verwenden:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: pgadmin
spec:
  replicas: 1
  template:
    spec:
      imagePullSecrets:
      - name: redhat-sso
      containers:
      - name: pgadmin
        image: registry.connect.redhat.com/crunchydata/crunchy-pgadmin4
```

## Verwenden der internen Registry :

```
$ oc patch configs.imageregistry.operator.openshift.io/cluster --patch '{"spec":{"defaultRoute":true}}' --type=merge  
(Administrator)
```

```
$ oc get route -n openshift-image-registry  
NAME          HOST/PORT  
default-route  default-route-openshift-image-registry.apps.eu46.prod.nextcle.com
```


Als Images-Repository wird der Namespace (Project) verwendet

default-route-openshift-image-registry.apps.eu46.prod.nextcle.com/<project>/<image>

```
$ skopeo list-tags docker://default-route-openshift-image-registry.apps.eu46.prod.nextcle.com/baseimages/rhel8  
{  
  "Repository": "default-route-openshift-image-registry.apps.eu46.prod.nextcle.com/baseimages/rhel8",  
  "Tags": [  
    "latest"  
  ]  
}
```

```
$ podman pull -creds $(oc whoami):$(oc whoami -t)  
docker://default-route-openshift-image-registry.apps.eu46.prod.nextcle.com/baseimages/rhel8  
Trying to pull docker://default-route-openshift-image-registry.apps.eu46.prod.nextcle.com/baseimages/rhel8...  
Getting image source signatures  
...
```

→ Zugriffsberechtigung !



```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: team:image-pullers
  namespace: do288-images
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:image-puller
subjects:
- apiGroup: rbac.authorization.k8s.io
  kind: Group
  name: system:serviceaccounts:do288-team1
- apiGroup: rbac.authorization.k8s.io
  kind: Group
  name: system:serviceaccounts:do288-team2
```

Verwenden einer externen Container Registry – Imagestream aktualisieren

```
$ oc import-image webserver --from=quay.io/do288/webserver --confirm
```

```
$ oc describe is webserver
```




```
Name: webserver
```

```
...
```

```
latest
```

```
  tagged from quay.io/do288/webserver
```

```
* quay.io/do288/webserver@sha256:1a618413d9a6cb45e37efc49a22cd08c5f702d6561483ed7dd2b38358e27fe10
```

TAG CHANGE		DATE/TIME	REVERT
 Mar 21, 2022			
	latest was moved to <b>SHA256</b> 2e43613a28b9 from <b>SHA256</b> 1a618413d9a6	Mon, Mar 21, 2022 9:16 PM	<a href="#">Revert to</a> <b>SHA256</b> 1a618413d9a6
	latest was moved to <b>SHA256</b> 1a618413d9a6 from <b>SHA256</b> cddd94b1691a	Mon, Mar 21, 2022 11:36 AM	<a href="#">Revert to</a> <b>SHA256</b> cddd94b1691a

```
$ oc tag quay.io/do288/webserver:latest webserver:latest
```

```
Tag webserver:latest set to quay.io/do288/webserver:latest
```

```
$ oc describe is webserver
```

```
...
```

```
latest
```


```
  tagged from quay.io/do288/webserver:latest
```

```
* quay.io/do288/webserver@sha256:2e43613a28b9614208adef64202646718e534b29e77328762c656c85793d37a9
  54 seconds ago
```

```
quay.io/do288/webserver@sha256:1a618413d9a6cb45e37efc49a22cd08c5f702d6561483ed7dd2b38358e27fe10
  7 minutes ago
```

Automatische Aktualisierung (15 min. Intervall) → `oc import-image webserver --from=quay.io/do288/webserver --scheduled`

## Image-Change



```
$ curl -XPOST https://api.eu46.prod.nextcle.com:6443/apis/build.openshift.io/v1/namespaces/images/buildconfigs/rhel8/webhooks/abcdefg/generic
```

```
$ oc get pods -w -n images
NAME          READY   STATUS
httpd-1-build 0/1     Completed
rhel8-1-build 0/1     Completed
rhel8-2-build 1/1     Running
httpd-2-build 0/1     Pending
rhel8-2-build 0/1     Completed
httpd-2-build 1/1     Running
httpd-2-build 0/1     Completed
```

```
$ oc get pods -w -n httpd-server
NAME          READY   STATUS
httpd-server-77b6fc6595-2487q 1/1     Running
httpd-server-cf6489bfd-qh2hr 0/1     Pending
httpd-server-cf6489bfd-qh2hr 0/1     ContainerCreating
httpd-server-cf6489bfd-qh2hr 1/1     Running
httpd-server-77b6fc6595-2487q 0/1     Terminating
```

```

apiVersion: build.openshift.io/v1
kind: BuildConfig
metadata:
  name: rhel8
  namespace: images
spec:
  source:
    dockerfile: |-
      FROM registry.access.redhat.com/ubi8/ubi:8.4
      ENV PACKAGES="lsf curl bind-utils"
      RUN dnf install -y --nodocs $PACKAGES && dnf clean all -y
    type: Dockerfile
  strategy:
    dockerStrategy: {}
    type: Docker
  output:
    to:
      name: rhel8:latest
      kind: ImageStreamTag
  successfulBuildsHistoryLimit: 1
  failedBuildsHistoryLimit: 1
  triggers:
  - type: Generic
    generic:
      secret: abcdefg
  - type: ConfigChange

```

## Build-Management:

```

oc start-build <name>
oc cancel-build <name>
oc set env bc/<name> BUILD_LOGLEVEL="4"


```

```

apiVersion: build.openshift.io/v1
kind: BuildConfig
metadata:
  name: httpd
  namespace: images
spec:
  source:
    type: Dockerfile
    dockerfile: |
      FROM xxxx
      RUN dnf install -y --nodocs httpd && dnf clean all -y
      ...
      EXPOSE 8080
      CMD /usr/sbin/httpd -DFOREGROUND
  strategy:
    type: Docker
    dockerStrategy:
      from:
        kind: ImageStreamTag
        namespace: images
        name: rhel8:latest
  successfulBuildsHistoryLimit: 1
  failedBuildsHistoryLimit: 1
  output:
    to:
      kind: ImageStreamTag
      name: httpd:latest
  triggers:
  - type: ImageChange

```





Source: binary | dockerfile | git | images

Strategy:

- source : Builder-Image enthält Tools und Logik zum Erstellen einer Anwendung (Source2Image)
- docker : Git-Repository mit Dockerfile

```

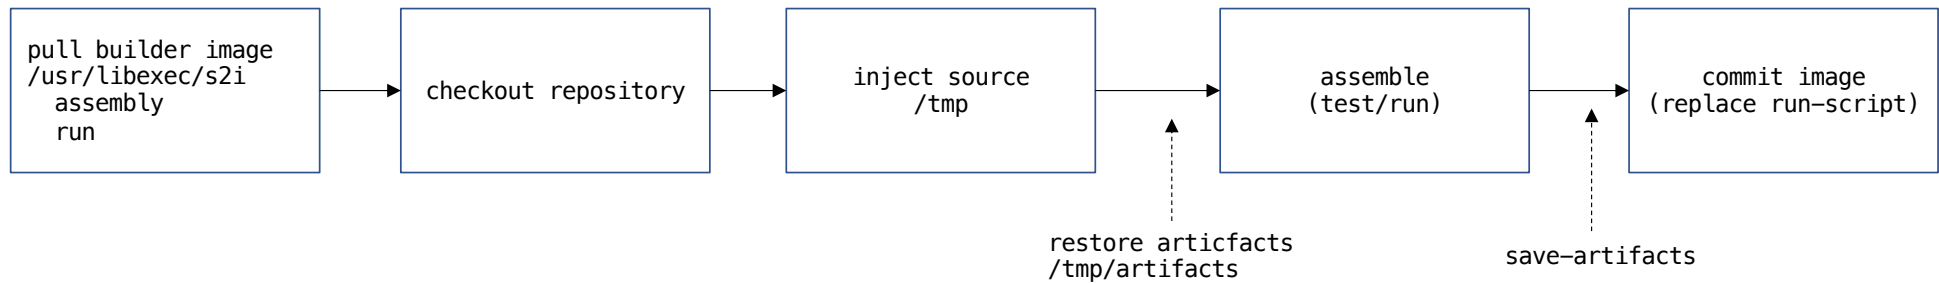
$ oc set triggers bc/sample --from-gitlab
buildconfig.build.openshift.io/sample triggers updated

$ oc describe bc/sample
...
Webhook GitHub:
  URL: https://api.eu46.prod.nextcle.com:6443/apis/build.openshift.io/v1/namespaces/dstraub-trigger/buildconfigs/sample/webhooks/<secret>/github
Webhook Generic:
  URL: https://api.eu46.prod.nextcle.com:6443/apis/build.openshift.io/v1/namespaces/dstraub-trigger/buildconfigs/sample/webhooks/<secret>/generic
Webhook GitLab:
  URL: https://api.eu46.prod.nextcle.com:6443/apis/build.openshift.io/v1/namespaces/dstraub-trigger/buildconfigs/sample/webhooks/<secret>/gitlab

$ $ oc get bc/sample -o json | jq '.spec.triggers'
[
  {
    "github": {
      "secret": "uU0xcyrsg4h58ThACUJj"
    },
    "type": "GitHub"
  },
  {
    "generic": {
      "secret": "Sfyo-MeJUbRGFS-3f0QX"
    },
    "type": "Generic"
  },
  ...
  {
    "gitlab": {
      "secret": "krhFoEzyiorD1UEZt_o5"
    },
    "type": "GitLab"
  }
]

$ curl -XPOST https://api.eu46.prod.nextcle.com:6443/apis/build.openshift.io/v1/namespaces/dstraub-trigger/buildconfigs/sample/webhooks/Sfyo-MeJUbRGFS-3f0QX/generic

```

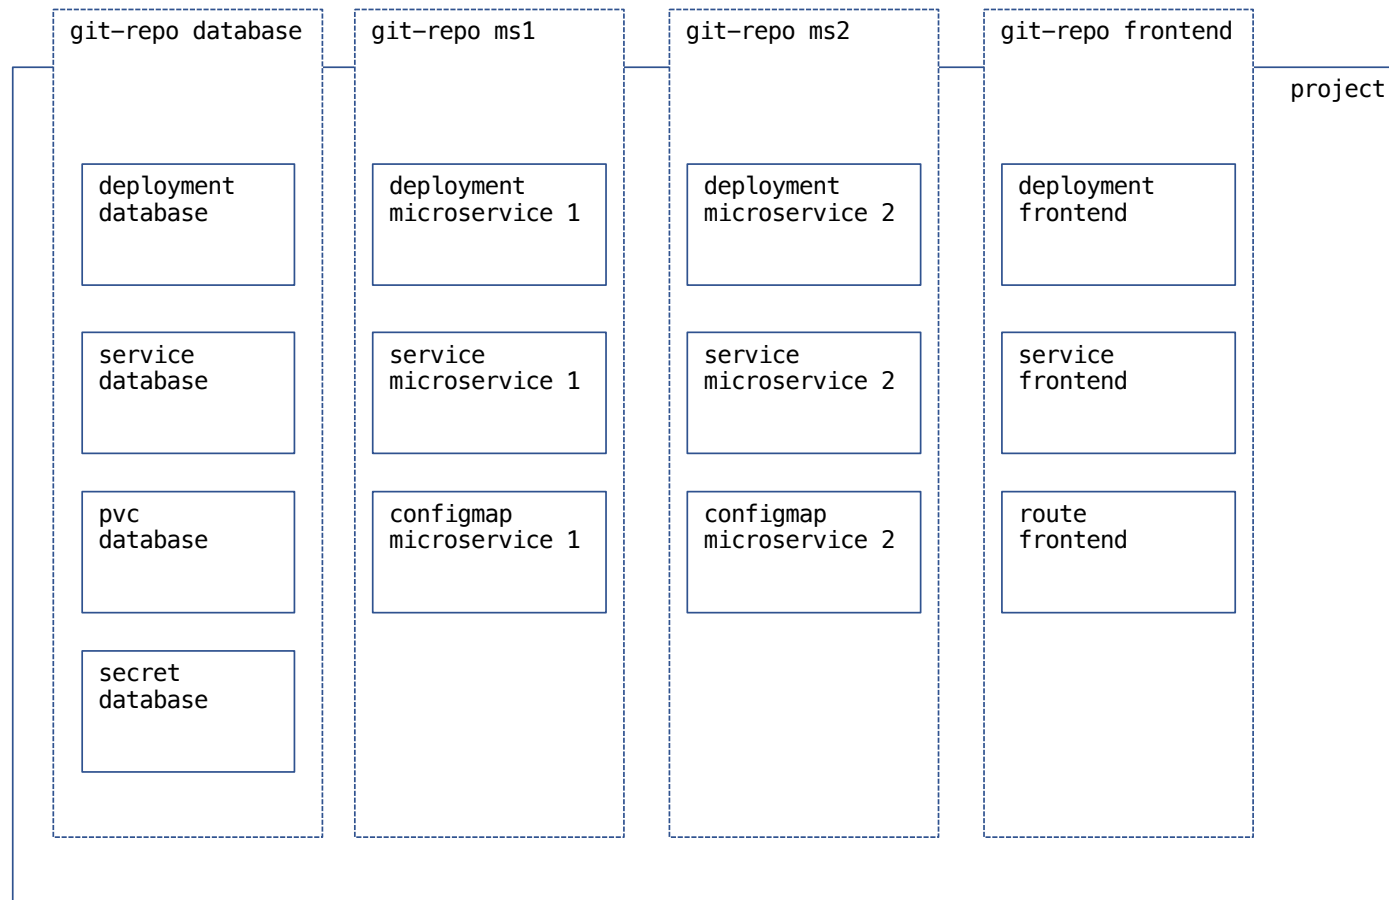


#### Build-Scripte:

- default in /usr/libexec/s2i
- assemble und run sind mandatory
- save-artifacts, usage, test/run sind optional
- können überschrieben werden im Git-Repo .s2i/bin  
(Wrapper um Original-Script oder komplett neues Script)

#### Incremental Builds:

- save-artifacts erstellt TAR
- wird vor dem Ausführen von assembly injected in /tmp/artifacts



## Template: parametrisierbare Liste von Resource-Definitionen

```
kind: Template
apiVersion: v1
metadata:
  name: rest-sample
objects:
- apiVersion: v1
  kind: Service
  metadata:
    name: ${APP_NAME}
  spec:
    selector:
      app.kubernetes.io/name: ${APP_NAME}
    ...
- apiVersion: apps/v1
  kind: Deployment
  metadata:
    name: ${APP_NAME}
  spec:
    template:
      spec:
        containers:
          - name: ${APP_NAME}
            image: ${IMAGE_NAME}
    ...
- apiVersion: v1
  kind: Route
  ...
parameters:
- description: Application Name
  name: APP_NAME
  required: true
- description: Image Name
  name: IMAGE_NAME
  required: true
...
```

```
oc process (TEMPLATE | -f FILENAME) -p APP_NAME=... | oc create -f -
oder bei installiertem Template ( oc create -f template.yml ) :
oc new-app <template-name>
```

```

apiVersion: template.openshift.io/v1
kind: Template
labels:
  app: php-sample
metadata:
  name: php-sample
  labels:
    samples.operator.openshift.io/managed: "true"
  app: php-sample
objects:
- apiVersion: v1
  kind: Service
  metadata:
    annotations:
      description: Exposes and load balances the application pods
  ...

```

Labels für alle Objekte

Labels nur für das Template

```

oc process -f template.yml -o yaml
apiVersion: v1
items:
- apiVersion: v1
  kind: Service
  metadata:
    labels:
      app: php-sample
  ...
- apiVersion: route.openshift.io/v1
  kind: Route
  metadata:
    labels:
      app: php-sample
  ...
- apiVersion: image.openshift.io/v1
  kind: ImageStream
  metadata:
    labels:
      app: php-sample
  ...

```

```

$ oc create -f template.yml
template.template.openshift.io/php-sample created

```

```

$ oc get template php-sample -o yaml
apiVersion: template.openshift.io/v1
kind: Template
labels:
  app: php-sample
metadata:
  labels:
    samples.operator.openshift.io/managed: "true"
  ...

```

\$ oc delete all -l app=php → alle vom Template erzeugten Objekte werden gelöscht

## Helm-Chart: Paket-Manager (Lifecycle + Template-Engine + Dependencies)

```
$ helm create sample
Creating sample

$ tree sample
sample
├── charts
├── Chart.yaml
├── templates
│   ├── deployment.yaml
│   ├── _helpers.tpl
│   ├── hpa.yaml
│   ├── ingress.yaml
│   ├── NOTES.txt
│   ├── serviceaccount.yaml
│   ├── service.yaml
│   └── tests
│       └── test-connection.yaml
└── values.yaml
```

## Helm-Chart: Paket-Manager (Lifecycle + Template-Engine + Dependencies)

```
Chart.yml
apiVersion: v1
name: sample
description: Sample Application
version: 1.0
appVersion: 1.0
dependencies:
- name: dep1
  version: ...
  repository: ...
```

```
values.yml
image:
  repository: quay.io/redhat.io/sample
  tag: '2'

service:
  port: 8080

env:
  ...

dep1.key: value
```

```
helm create
helm dependency update
helm install / upgrade / rollback / uninstall

helm template (lokales processing)
```

### Templates:

```
deployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ APP_NAME }}
spec:
  template:
    selector:
      matchLabels:
        {{- include "sample.selectorLabels" . | nindent 6 }}
    spec:
      containers:
        - image: ${.Values.image.repository}: ${.Values.image.tag}
      ...
```

### Go-Templates:

```
_helpers.tpl
{{- define "sample.selectorLabels" -}}
app.kubernetes.io/name: {{ include "sample.name" . }}
app.kubernetes.io/instance: {{ .Release.Name }}
{{- end -}}
...
```



## Kustomize: generieren/transformieren von Ressourcen (Manifeste mit minimalen Meta-Daten)

```
kind: Kustomization                                kustomization.yml
apiVersion: kustomize.config.k8s.io/v1beta1

namespace: sample

resources:
- deployment.yml
- service.yml
- route.yml
- http://...    -> kustomize.yml in Git/Web-Repository

images:
- name: sample
  newName: registry/sample
  newTag: '5'

commonLabels:
  app.kubernetes.io/instance: sample

configMapGenerator:
- name: rest-sample
  literals:
  - LAUNCH_JBOSS_IN_BACKGROUND=1
...
```

resources → <https://github.com/hashicorp/go-getter#url-format>

```
apiVersion: apps/v1
metadata:
  name: rest-sample
spec:
  replicas: 1
  template:
    spec:
      containers:
      - name: sample
        image: sample
```

```
$ oc kustomize <kustom-dir>
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/instance: rest-sample
  name: rest-sample
  namespace: sample
spec:
  replicas: 1
  selector:
    matchLabels:
      app.kubernetes.io/instance: sample
  template:
    containers:
      image: registry/sample:5
...

$ oc apply -k .
```

Kustomize Overlays : erzeugen unterschiedlicher Varianten von einer Basis-Vorlage

```
                                base/kustomization.yml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- deployment.yml
- service.yml
- route.yml
```

```
                                overlays/test/kustomization.yml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- ../../base

namespace: test

images:
- name: sample
  newName: registry/sample
  newTag: '3-SNAPSHOT'
```

```
                                overlays/production/kustomization.yml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- ../../base

namespace: production

images:
- name: sample
  newName: registry/sample
  newTag: '5'
```

```
$ oc apply -k overlays/test
service/sample configured
deployment.apps/sample configured
route.route.openshift.io/sample configured

$ oc apply -k overlays/production
...
```

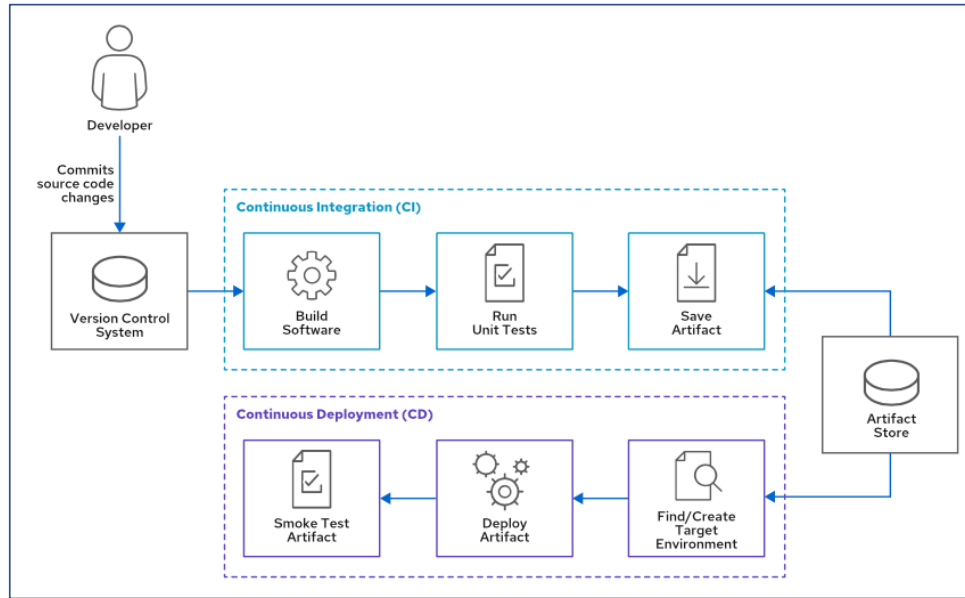
[https://kubectldocs.kubernetes.io/guides/extending\\_kustomize/exec\\_krm\\_functions](https://kubectldocs.kubernetes.io/guides/extending_kustomize/exec_krm_functions)

```
$ grep -A3 images kustomization.yml
images:
- name: webserver
  newName: quay.io/danielstraub/webserver
  newTag: "1.0"

$ kustomize edit set image webserver=quay.io/danielstraub/webserver:2.0
$ grep -A3 images kustomization.yml
images:
- name: webserver
  newName: quay.io/danielstraub/webserver
  newTag: "2.0"

$ oc apply -k .
configmap/webserver-kt5mdg45d2 unchanged
service/webserver unchanged
deployment.apps/webserver configured
route.route.openshift.io/webserver unchanged

$ curl https://stage-prod.apps.eu46a.prod.ole.redhat.com
Hello, D0288
Version 2.0
```



Continuous Integration  
Continuous Delivery

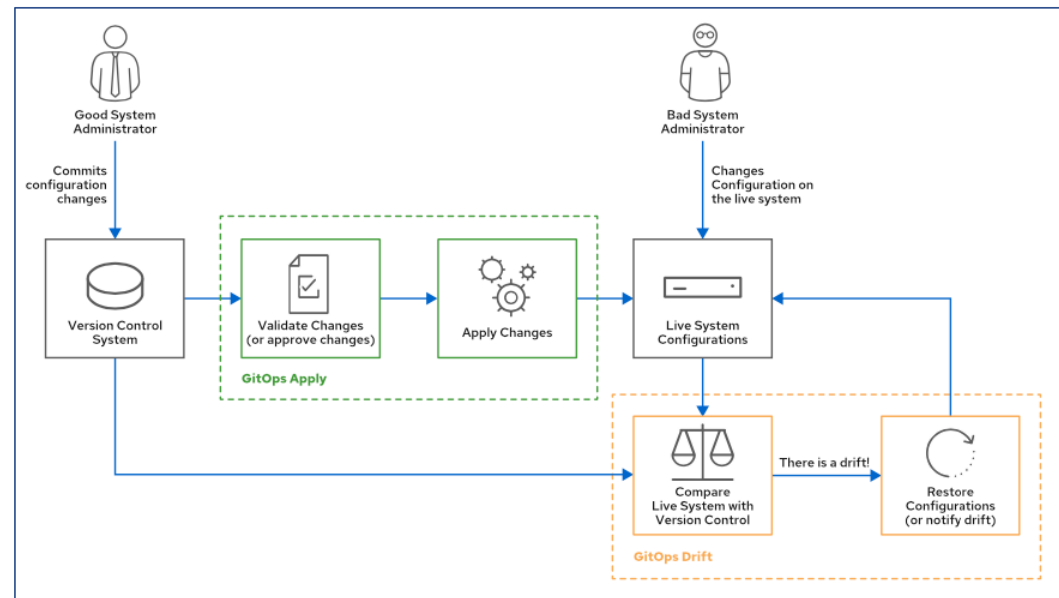
→ Developer  
→ running application

Jenkins, CruiseControl, TeamCity, GitLab ...  
Kubernetes native (Tekton, Argo CD, ...)

## GitOps Workflow

→ Administrators  
→ live System

Ansible, Puppet, Terraform ...  
ArgoCD, FluxCD, JenkinsX



## GitOps – Workflow mit Pipelines:

- Apply Pipeline:
  - validate : `oc apply --validate --dry-run [ folder/files from Git ]`
  - apply : `oc apply`
- Drift Pipeline:
  - diff : `oc diff [ folder/files from Git ]`
  - optional/restore: `oc apply`


## GitOps – Workflow mit ArgoCD / FluxCD:

Ableich Ist-Zustand (Cluster) mit Kustomize/Helm-Definitionen im Git

Benachrichtigungen, manueller/automatische Synchronisation bei Abweichungen

apps calibre	ssh://git@gitea.apps:10022/ds/calibre.git/overlays/production in-cluster/apps	HEAD	♥ Healthy ✓ Synced	⋮
apps pgadmin	ssh://git@gitea.apps:10022/ds/pgadmin.git/overlays/production in-cluster/apps	HEAD	♥ Healthy ✓ Synced	⋮
apps postgres	ssh://git@gitea.apps:10022/ds/postgres.git/overlays/production in-cluster/database	HEAD	♥ Healthy ✓ Synced	⋮
apps rest-sample	ssh://git@gitea.apps:10022/ds/rest-sample.git/overlays/production in-cluster/sample	HEAD	♥ Healthy ⚠ OutOfSync	⋮

## Red Hat OpenShift GitOps - Operator



ArgoCD = Openshift GitOps

← → ↻ 🏠 <https://openshift-gitops-server-openshift-gitops.apps.ocp4.example.com/applications/stage-prod?view=tree&re> ☆

Applications / stage-prod APPLICATION DETAILS

APP DETAILS APP DIFF SYNC SYNC STATUS HISTORY AND ROLLBACK DELETE REFRESH

APP HEALTH **Healthy**

CURRENT SYNC STATUS **OutOfSync** From **HEAD (f46fb2b)** [MORE](#)

LAST SYNC RESULT **Sync OK** To **5852439** [MORE](#)

Succeeded 14 hours ago (Wed Jun 29 2022 13:13:26 GMT-0400)  
Author: Daniel Straub <ds@ctrlaltdel.de> - update sync policy  
Comment:

Succeeded 14 hours ago (Wed Jun 29 2022 13:13:26 GMT-0400)  
Author: Daniel Straub <ds@ctrlaltdel.de> - Update kustomization.yml  
Comment:

stage-prod

webserver-kt5mdg45d2 (cm) 14 hours

webserver (svc) 14 hours

webserver (deploy) 14 hours

webserver (route) 14 hours

webserver (ep) 14 hours

webserver-b88c4 (ES endpointslice) 14 hours

webserver-567899c8fd (rs) 14 hours rev:1

webserver-7676d986b9 (rs) 14 hours rev:2

webserver (pod) 14 hours

```
$ oc get pods -w
```

NAME	READY	STATUS	RESTARTS	AGE
famousapp-famouschart-65744d4c8b-4zqhn	0/1	Running	0	10s
famousapp-mariadb-0	0/1	ContainerCreating	0	10s
famousapp-mariadb-0	0/1	Running	0	11s
famousapp-famouschart-65744d4c8b-4zqhn	0/1	Running	1	33s
famousapp-famouschart-65744d4c8b-4zqhn	0/1	Error	1	34s
famousapp-famouschart-65744d4c8b-4zqhn	0/1	Running	2	35s
famousapp-famouschart-65744d4c8b-4zqhn	0/1	Error	2	36s
famousapp-famouschart-65744d4c8b-4zqhn	0/1	CrashLoopBackOff	2	37s
famousapp-mariadb-0	1/1	Running	0	48s
famousapp-famouschart-65744d4c8b-4zqhn	0/1	Running	3	56s
famousapp-famouschart-65744d4c8b-4zqhn	1/1	Running	3	62s

```

metadata:
  name: famousapp-mariadb:
...

livenessProbe:
  exec:
    command:
      - /bin/bash
      - -ec
      - |
        password_aux="${MARIADB_ROOT_PASSWORD:-}"
        if [[ -f "${MARIADB_ROOT_PASSWORD_FILE:-}" ]]
        password_aux=$(cat "${MARIADB_ROOT_PASSWORD_FILE:-}")
        fi
        mysqladmin status -uroot -p"${password_aux}"

```

```

metadata:
  name: famousapp-famouschart
...
livenessProbe:
  initialDelaySeconds: 30
  httpGet:
    path: /
...
readinessProbe:
  failureThreshold: 3
  httpGet:
    path: /
    port: http
    scheme: HTTP
    periodSeconds: 10
    successThreshold: 1
    timeoutSeconds: 1

```



## Liveness / Readiness / Startup Probes

- liveness : Container wird bei negativen Ergebnis neu gestartet `.spec.containers.livenessProbe`
- readiness: Route/Service wird aktiviert/deaktiviert `.spec.containers.readinessProbe`
- startup: liveness/readiness sind deaktiviert bis startup positiv ist `.spec.containers.startupProbe`  
Container wird bei neg. Startup-Probe sofort beendet

### Probes:

```
exec:
  command:
  - cat
  - /tmp/ready
  initialDelaySeconds: 5
  periodSeconds: 5
```

```
httpGet:
  path: /healthz
  port: healthz-port
  schema: https
  httpHeaders: ...
  failureThreshold: 1
  periodSeconds: 10

200 <= status < 400
```

```
tcpSocket:
  port: 5432
  initialDelaySeconds: 15
  periodSeconds: 20
```

- initialDelaySeconds: Zeitdauer bis zur ersten liveness/readiness Probe
- periodSeconds: Intervall zur Ausführung der Proben (default 10 sec)
- timeoutSeconds: max. Timeout bei einer Probe (default 1 sec)
- successThreshold: Schwellwert ab wann aufeinanderfolgende positive Proben als Erfolg gewertet werden (default 1)
- failureThreshold: Schwellwert ab wann aufeinanderfolgende negative Proben als Ausfall gewertet werden (default 3)

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: webserver
spec:
  ...
  template:
    spec:
      containers:
      - name: webserver
        image: webserver
        imagePullPolicy: Always
        ports:
        - name: http
          containerPort: 8080
          protocol: TCP
        readinessProbe:
          failureThreshold: 3
          httpGet:
            path: /healthz
            port: http
            scheme: HTTP
          periodSeconds: 10
          successThreshold: 1
          timeoutSeconds: 1
      ...
```

```
server {
    listen 8080 default_server;
    server_name _;
    location / {
        root    /usr/share/nginx/html;
        index   index.html index.htm;
    }

    location /healthz {
        access_log off;
        return 200;
    }
}
```

nginx.conf

## DeploymentConfig | Deployment

```
kind: DeploymentConfig
metadata:
  name: ...
spec:
  replicas: 1
  selector:
    app: ...
  template:
    metadata:
      ...
    spec:
      strategy:
        rollingParams:
          pre:
          mid:
          post:
      containers:
      - name: <container_name>
        image: image-registry.openshift-image-registry.svc:5000/<name_space>/<image>:@sha256:xxxx
        imagePullPolicy: IfNotPresent
        ...
      triggers:
      - type: ConfigChange
      - type: ImageChange
        imageChangeParams:
          containerNames:
          - <container_name>
          from:
            name: <image_stream>:<image_tag>
```

Automatisches Redeployment bei Konfigurations-Änderungen oder neues Image im verknüpften Imagestream

```
kind: Deployment
metadata:
  name: <name>
  annotations:
    image.openshift.io/triggers: '{"from": {"kind": "ImageStreamTag", "name": "<image_stream>:<image_tag>"}, "fieldPath": "spec.template.spec.containers[0].image" }'
spec:
```

## Deployment-Strategien

- Rolling Updates : Pods werden der Reihe nach aktualisiert
- Recreate: existierende Pods werden beendet und neue gestartet

## DeploymentConfig:

- Pre/Mid/Post – Lifecycle Hooks

## Beenden eines Pods:

- SIGTERM: Pod soll keine neuen Verbindungen annehmen und bestehenden Aktionen beenden
- SIGKILL: nach `terminationGracePeriodSeconds` (30s) wird der Pod beendet


```
kind: Deployment
metadata:
  name: ...
spec:
  revisionHistoryLimit: 3   (default: 10)
  replicas: 4
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1           ← max. 5 Pods aktiv
      maxUnavailable: 1
  ...
  template:
    spec:
      containers:
      - ...
      terminationGracePeriodSeconds: 30
```

```
oc rollout SUBCOMMAND (DEPLOYMENTCONFIG | DEPLOYMENT)
```

cancel	Cancel the in-progress deployment
history	View rollout history
latest	Start a new rollout for deployment config with latest state
pause	Mark the provided resource as paused
restart	Restart a resource
resume	Resume a paused resource
retry	Retry the latest failed rollout
status	Show the status of the rollout
undo	Undo a previous rollout

```
oc rollback (DEPLOYMENTCONFIG | DEPLOYMENT) [--to-version=]
```

## N-1 Abwärtskompatibilität bei Rolling-Update:



## A/B Deployment Strategy:

```
apiVersion: v1
kind: Service
metadata:
  name: service-a
spec:
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: http
  selector:
    app.kubernetes.io/instance: deployment-a
```

```
apiVersion: v1
kind: Service
metadata:
  name: service-b
spec:
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: http
  selector:
    app.kubernetes.io/instance: deployment-b
```


```
kind: Route
metadata:
  name: <name>
spec:
  host: <host>
  to:
    kind: Service
    name: service-a
    weight: 50
  alternateBackends:
    - kind: Service
      name: service-b
      weight: 200
```

Service :

- ClusterIP
- NodePort
- Loadbalancer (AWS,Azure, MetalLB ...)
- ExternalName

StatefulSets → Headless Service

```
apiVersion: v1
kind: Service
metadata:
  name: webserver
spec:
  type: ClusterIP
  selector:
    app.kubernetes.io/instance: webserver
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: http
```



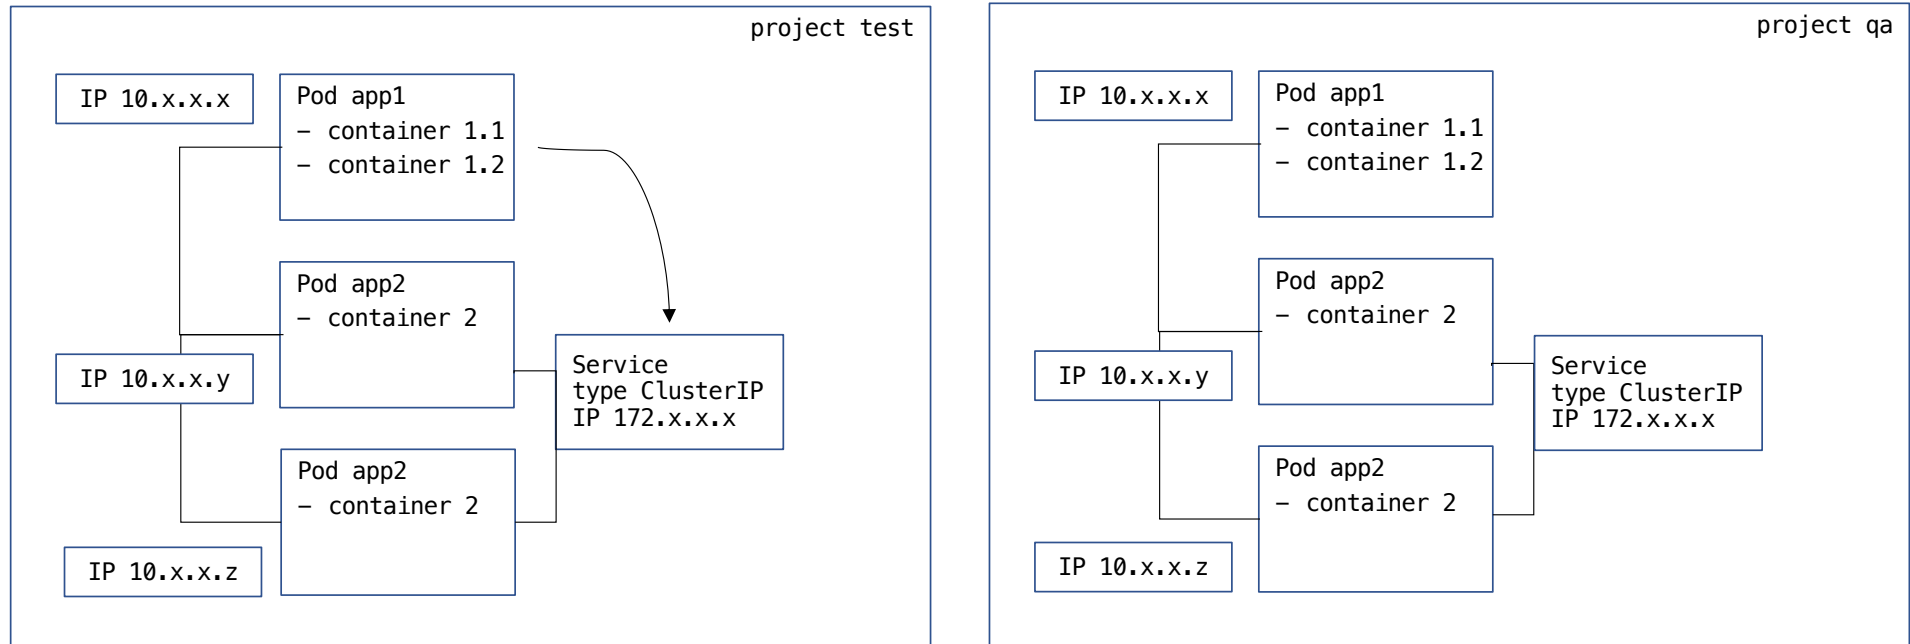


## Pod | Service | Route

```
apiVersion: v1
kind: Pod
metadata:
  name: webserver
  labels:
    app.kubernetes.io/instance: httpd
spec:
  containers:
  - name: httpd
    image: ...
    ports:
    - name: http
      containerPort: 8080
    - name: https
      containerPort: 8443
```

```
apiVersion: v1
kind: Service
metadata:
  name: webserver
spec:
  selector:
    app.kubernetes.io/instance: httpd
  ports:
  - name: http
    port: 80
    protocol: TCP
    targetPort: http
  - name: https
    port: 443
    protocol: TCP
    targetPort: https
```

```
apiVersion: route.openshift.io/v1
kind: Route
metadata:
  name: webserver-secure
spec:
  host: webserver.apps....
  to:
    kind: Service
    name: webserver
  port:
    target-port: https
```



DNS:

A: <service>.test.svc.cluster.local

SVC: \_443.\_tcp.https.<service>.test.svc.cluster.local

/etc/resolv.conf:

search test.svc.cluster.local svc.cluster.local ...

ndots 5

DNS:

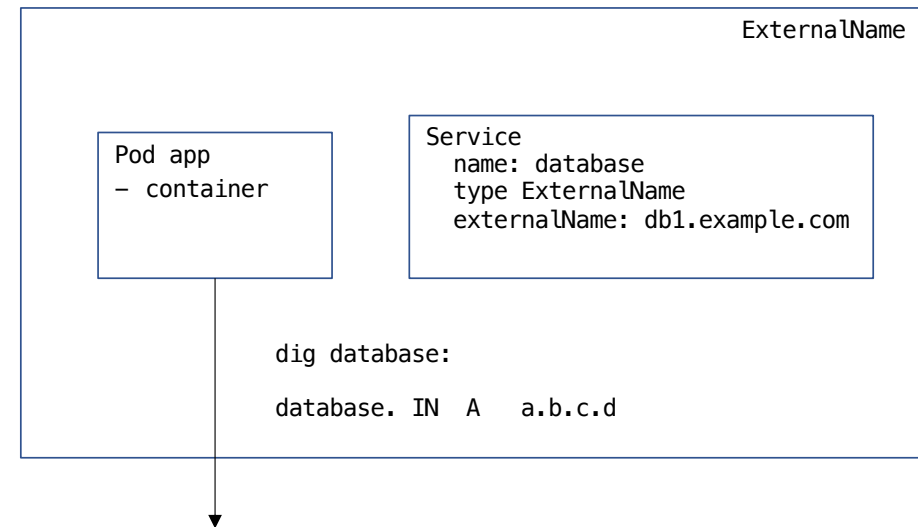
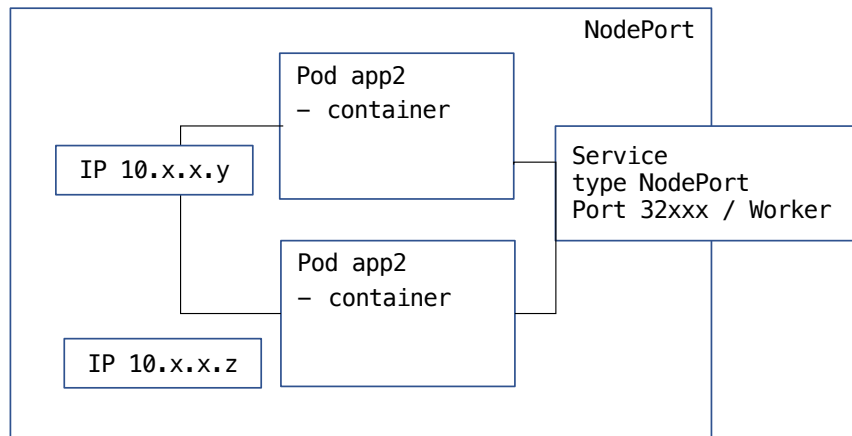
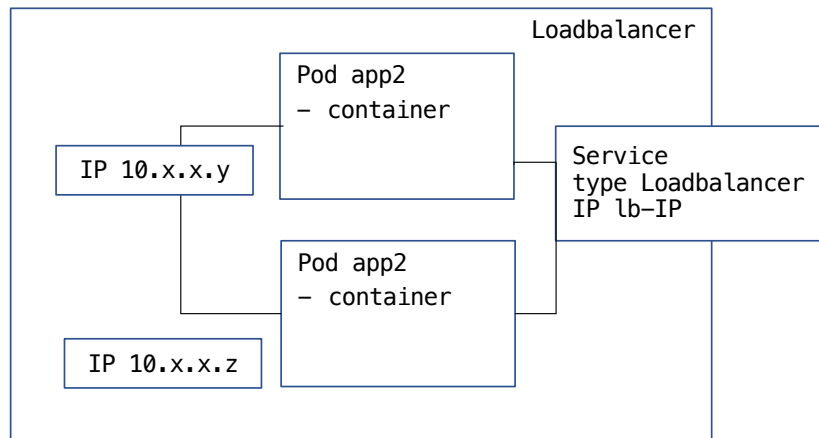
A: <service>.qa.svc.cluster.local

SVC: \_443.\_tcp.https.<service>.qa.svc.cluster.local

/etc/resolv.conf:

search qa.svc.cluster.local svc.cluster.local ...


→ einfacher DNS-Lookup nach <service> in jedem Projekt



db1.example.com  
a.b.c.d

```
kind: StatefulSet
metadata:
  name: app
spec:
  ...
```

#### Headless Service



Service  
**clusterIP: None**

→ kein Loadbalancing, A-Record für jeden Host

```
dig app-1:
IN A 10.x.x.y
```

```
dig app-2
IN A 10.x.x.y
```

## Services für externe Datenbank:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
database	ExternalName	<none>	dsbox.de	<none>

```
apiVersion: v1
kind: Service
metadata:
  name: database
  namespace: sample
spec:
  type: ExternalName
  externalName: dsbox.de
```

```
# oc exec rest-sample-59fc6bf5b6-9dchd -- sh -c 'psql postgresql://daniel:12345678@database/postgres -c "\conninfo"'
You are connected to database "postgres" as user "daniel" on host "database-ext" at port "5432".
```

### DNS-Gotcha:

```
# oc exec rest-sample-59fc6bf5b6-9dchd -- sh -c 'nslookup dsbox.de'
Name:      dsbox.de.<wildcard.domain>
Address: 5.9.70.75
```

```
# oc exec rest-sample-59fc6bf5b6-9dchd -- sh -c 'nslookup dsbox.de.'
Name:      dsbox.de
Address: 176.9.155.194
```

<https://imgtfy.app/?q=options+ndots%3A5>

## Services für Database-Pod:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
database	ClusterIP	172.30.156.203	<none>	5432/TCP
database-node	NodePort	172.30.160.12	<none>	5432:30001/TCP
database-ip	ClusterIP	172.30.31.229	10.0.0.21,10.0.0.22	5432/TCP

```
apiVersion: v1
kind: Service
metadata:
  name: database
  namespace: sample
spec:
  selector:
    app.kubernetes.io/name: database
  type: ClusterIP
  ports:
    - name: database
      protocol: TCP
      port: 5432
      targetPort: 5432
```

```
apiVersion: v1
kind: Service
metadata:
  name: database-node
  namespace: sample
spec:
  selector:
    app.kubernetes.io/name: database
  type: NodePort
  ports:
    - name: database
      protocol: TCP
      port: 5432
      targetPort: 5432
      nodePort: 30001 ← Range 30000-32000
```

```
apiVersion: v1
kind: Service
metadata:
  name: database-ip
  namespace: sample
spec:
  selector:
    app.kubernetes.io/name: database
  ports:
    - name: database
      protocol: TCP
      port: 5432
      targetPort: 5432
  externalIPs:
    - 10.0.0.21
    - 10.0.0.22
```

```
# oc exec rest-sample-59fc6bf5b6-9dchd -- sh -c 'psql postgresql://daniel:12345678@database/postgres -c "\conninfo"'
You are connected to database "postgres" as user "daniel" on host "database" at port "5432".
```

```
# psql postgresql://daniel:12345678@worker01:30001/postgres -c "\conninfo"
You are connected to database "postgres" as user "daniel" on host "worker01" at port "30001"
```

```
# psql postgresql://daniel:12345678@10.0.0.21/postgres -c "\conninfo"
You are connected to database "postgres" as user "daniel" on host "10.0.0.21" at port "5432"
```

```
# psql postgresql://daniel:12345678@worker02/postgres -c "\conninfo"
You are connected to database "postgres" as user "daniel" on host "worker02" at port "5432"
```

## Pod – Scheduling

### 1. Filter

- Node-Selector für Labels  
<https://kubernetes.io/docs/reference/labels-annotations-taints>
- Toleration für Taints  
<https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration>

```
apiVersion: v1
kind: Pod
metadata:
  ...
spec:
  containers:
  - name: nginx
    nodeSelector:
      disktype: ssd
  tolerations:
  - key: class
    value: do280
    operator: "Equal"
    effect: "NoSchedule"
```

```
apiVersion: v1
kind: Node
metadata:
  labels:
    disktype: ssd
spec:
  taints:
  - key: class
    value: do280
    effect: NoSchedule
```

## Pod – Scheduling

### 2. Scoring

#### Affinity/Anti-Affinity-Rules

```
apiVersion: v1
kind: Pod
metadata:
  name: with-node-affinity
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: kubernetes.io/os
                operator: In
                values:
                  - linux
      preferredDuringSchedulingIgnoredDuringExecution:
        - weight: 1
          preference:
            matchExpressions:
              - key: another-node-label-key
                operator: In
                values:
                  - another-node-label-value
  containers:
    - name: with-node-affinity
      image: ...
```

...DuringScheduling: während des Scheduling

IgnoredDuringExecution: Pod wird weiter ausgeführt,  
auch wenn sich nach dem Scheduling Node-Labels ändern

<https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node>

<https://www.cncf.io/blog/2021/07/27/advanced-kubernetes-pod-to-node-scheduling>



## Pod Verteilung auf unterschiedliche Nodes:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  ...
spec:
  selector:
    matchLabels:
      app: store
  replicas: 3
  template:
    metadata:
      labels:
        app: store
    spec:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: app
                    operator: In
                    values:
                      - store
              topologyKey: "kubernetes.io/hostname"
      containers:
        - ...
```

gleiche Label und gleicher Hostname → AntiAffinity  
Pods werden auf unterschiedlichen Nodes verteilt

[https://docs.openshift.com/container-platform/4.10/nodes/scheduling/nodes-scheduler-pod-affinity.html#nodes-scheduler-pod-affinity-example-antiaffinity\\_nodes-scheduler-pod-affinity](https://docs.openshift.com/container-platform/4.10/nodes/scheduling/nodes-scheduler-pod-affinity.html#nodes-scheduler-pod-affinity-example-antiaffinity_nodes-scheduler-pod-affinity)

(compute) Resources :

- Memory: number of bytes (quantity suffixes: E, P, T, G, M, k | Ei, Pi, Ti, Gi, Mi, Ki)
- CPU : millicores (m)

*millicores* are the fractions of *time* of a single CPU (not the fraction of number of CPUs).

*Cgroups*, and hence Docker, and hence Kubernetes, doesn't restrict CPU usage by assigning cores to processes (like VMs do), instead it restricts CPU usage by restricting the amount of time (quota over period) the process can run on each CPU (with each CPU taking up to 1000mcpus worth of allowed time).

<https://stackoverflow.com/questions/61851751/multi-threading-with-millicores-in-kubernetes>

```
apiVersion: v1
kind: Pod
metadata:
...
spec:
  containers:
  - name: <name>
    resources:
      requests:
        memory: 64Mi
        cpu: 100m
      limits:
        memory: 128Mi
        cpu: 200m
```

Scheduling

Execution  
(cgroups)

```
$ oc describe node master01
...
Allocatable:
  cpu:          3500m
  memory:       15268156Ki
Non-terminated Pods: (60 in total)
   CPU Requests CPU Limits Memory Requests Memory Limits
...
Allocated resources:
Resource          Requests          Limits
-----
cpu               2397m (68%)      0 (0%)
memory           9347Mi (62%)     512Mi (3%)
```